

Lesson 1

Week 1

Lesson 1 - Introduction to Blockchain and Layer 1

Lesson 2 - Why Scalability

Lesson 3 - Introduction to Layer 2

Lesson 4 - Maths and Cryptography

Today's topics

- Decentralised Systems
 - Blockchain timeline
 - Layer 1 Theory
 - Consensus on Ethereum
-

Course Introduction

Course Plan

Week 1 Background / Scalability

- Intro to L1
- Why Scalability
- Intro to L2s
- Maths and Cryptography

Week 2 -

- Criteria for Understanding and Analysing L2s
- Agnostic L2 Transaction Lifecycle
- Optimistic Rollups v ZK Rollups
- What's next in L2s (1) - Decentralized Sequencers

Week 3

- What's next in L2s (2) - Privacy, L3s/Hyperchains
- Privacy in L2s
- What are ZK EVMs 1 - overview
- What are ZK EVMs 2 - universal circuits/circuit compiler

Week 4

- What are ZK EVMs 3 - proving system: zk algorithm/ASIC optimisation
- zkEVM security
- Overview of Proving Systems
- SNARK implementation

Week 5

- STARK implementation
- Plonk part 1 / Linea
- Plonk part 2 / Boojum
- Advanced Plonk / Scroll

- Formal Verification
- Risc Zero
- ZK Machine Learning
- Review

Practical Details

The format will usually be 45 mins of theory followed by 45 mins practical

You do not need to submit the homeworks

You can work in teams in breakout rooms

We use Sli.do to provide Q&A and polls : [link](#)

About Extropy

ABOUT US

Extropy.io was founded 2015 by Laurence Kirk in Oxford to provide consultancy services in Distributed Ledger Technology. Laurence is also the founder of the Oxford Blockchain Society.

**INNOVATE.
QUALITY.
CUTTING EDGE.**



CONTACT US

Oxford Centre for Innovation, New Road, Oxford, OX1 1BY, UK
www.extropy.io
+44 (0)1865 261 424

Providing Blockchain solutions
DApp development and customised blockchains
Security Audits



EXTROPY.IO
CONSULTANCY IN DISTRIBUTED LEDGER TECHNOLOGY

Free Developer Workshops

- Basic
- Enterprise
- Advanced EVM
- Zero Knowledge Proofs

Business Workshops

Website :

<https://extropy.io>

Email :

info@extropy.io

Twitter : [@extropy](#)

Social Media : @Extropy

<https://discord.gg/XzZS3FBx>

General Resources

[Zero Knowledge Podcast](#)

Decentralised Systems

Problems with centralised systems

Monetary System

- Bank closure / insufficient capital reserves
- greek debt crisis in 2015 ? banks closed and people lost savings, insurance schemes meant nothing, lead to an increase in Bitcoin use in Greece
- Availability of banks
- Inflation - money supply controlled by central authority
- Merchant accounts may be shut down
- Control of money for political reasons - wikileaks funding shutdown

There are layers of access control built into our banking systems to prevent fraudulent transactions, effectively security is achieved by closing the network.

Goals of decentralisation

- Participation
- Diversity
- Conflict resolution
- Flexibility
- Moving power to the edge (user)

Blockchain Timeline

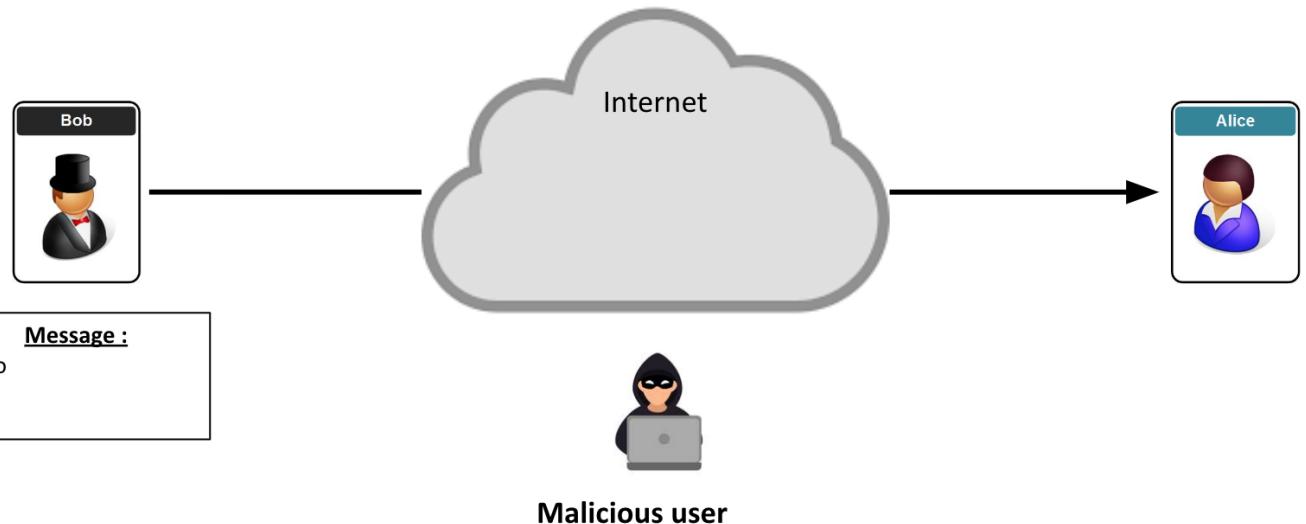
1970s

Problem = Security !

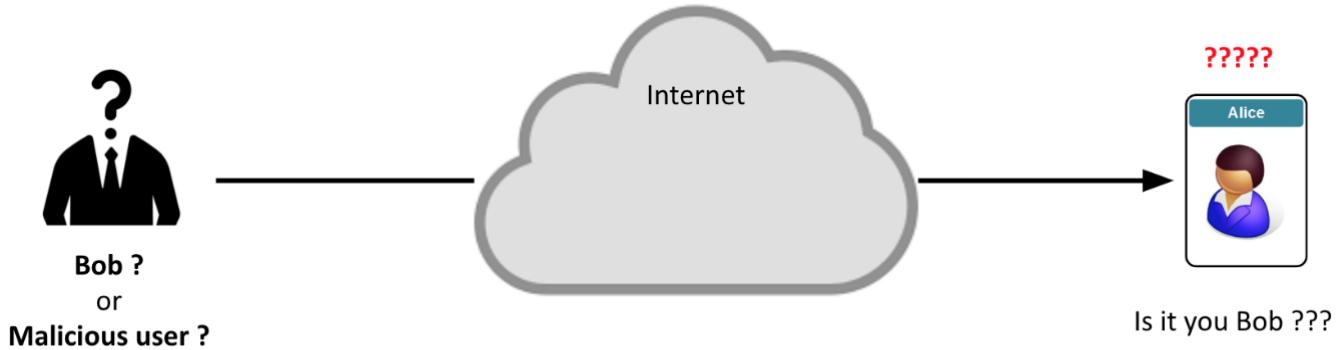
- How do I ensure that my message has not been modified ?
- How do I ensure that the message comes from a legitimate person ?

Secure Communication over Insecure Channel

Problem 1 : How do I ensure that my message has not been modified ?



Problem 2 : How do I ensure that the message comes from a legitimate person ?



pre 1970s solution : Symmetric Cryptography !

- Alice and Bob share the same key.
- One key for both encryption and decryption of messages

But what about key management ?

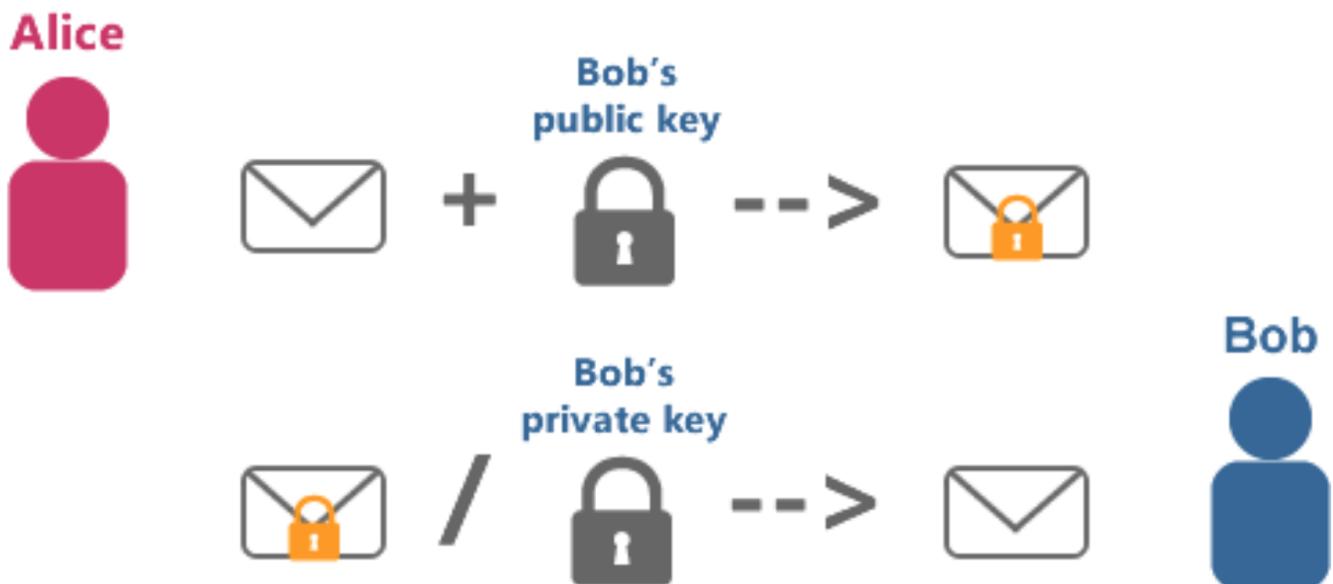
Can Alice and Bob share a key

- Without meeting
- Across a potentially hostile network

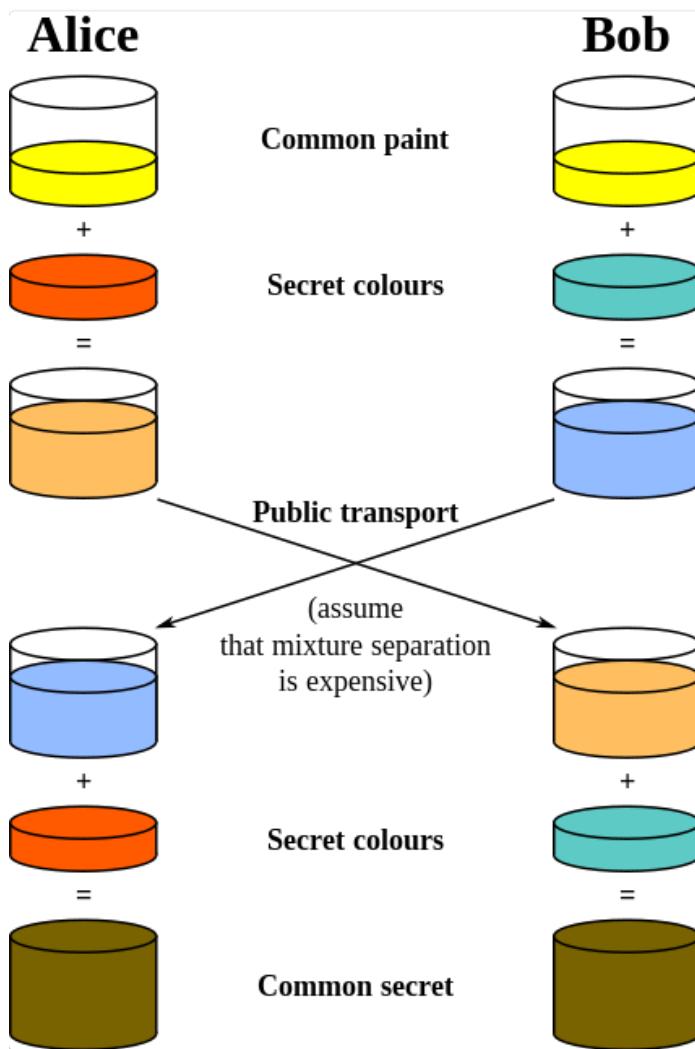
In the 1970s there was a major step forward with the invention of asymmetric cryptography.

Here 2 paired keys are used, one for encryption and one for decryption.

Cryptography - Asymmetric Keys



A method was also designed to allow the production of a shared secret across an insecure channel, via Diffie-Helman Key Exchange. An analogy of the process is provided by thinking of colours being mixed. See [Guide] (<https://github.com/archit-p/simplest-oblivious-transfer>)



Public Key Encryption solves :



Problem 1 : Key Management

- If I use a 3rd party to share my key, do I trust him ?

Problem 2 : Integrity

- How do I ensure that my message has not been modified ?

Problem 3 : Authenticity

- How do I ensure that the ~~message~~ comes from a legitimate person ?



This final problem was solved by the invention of digital signatures.

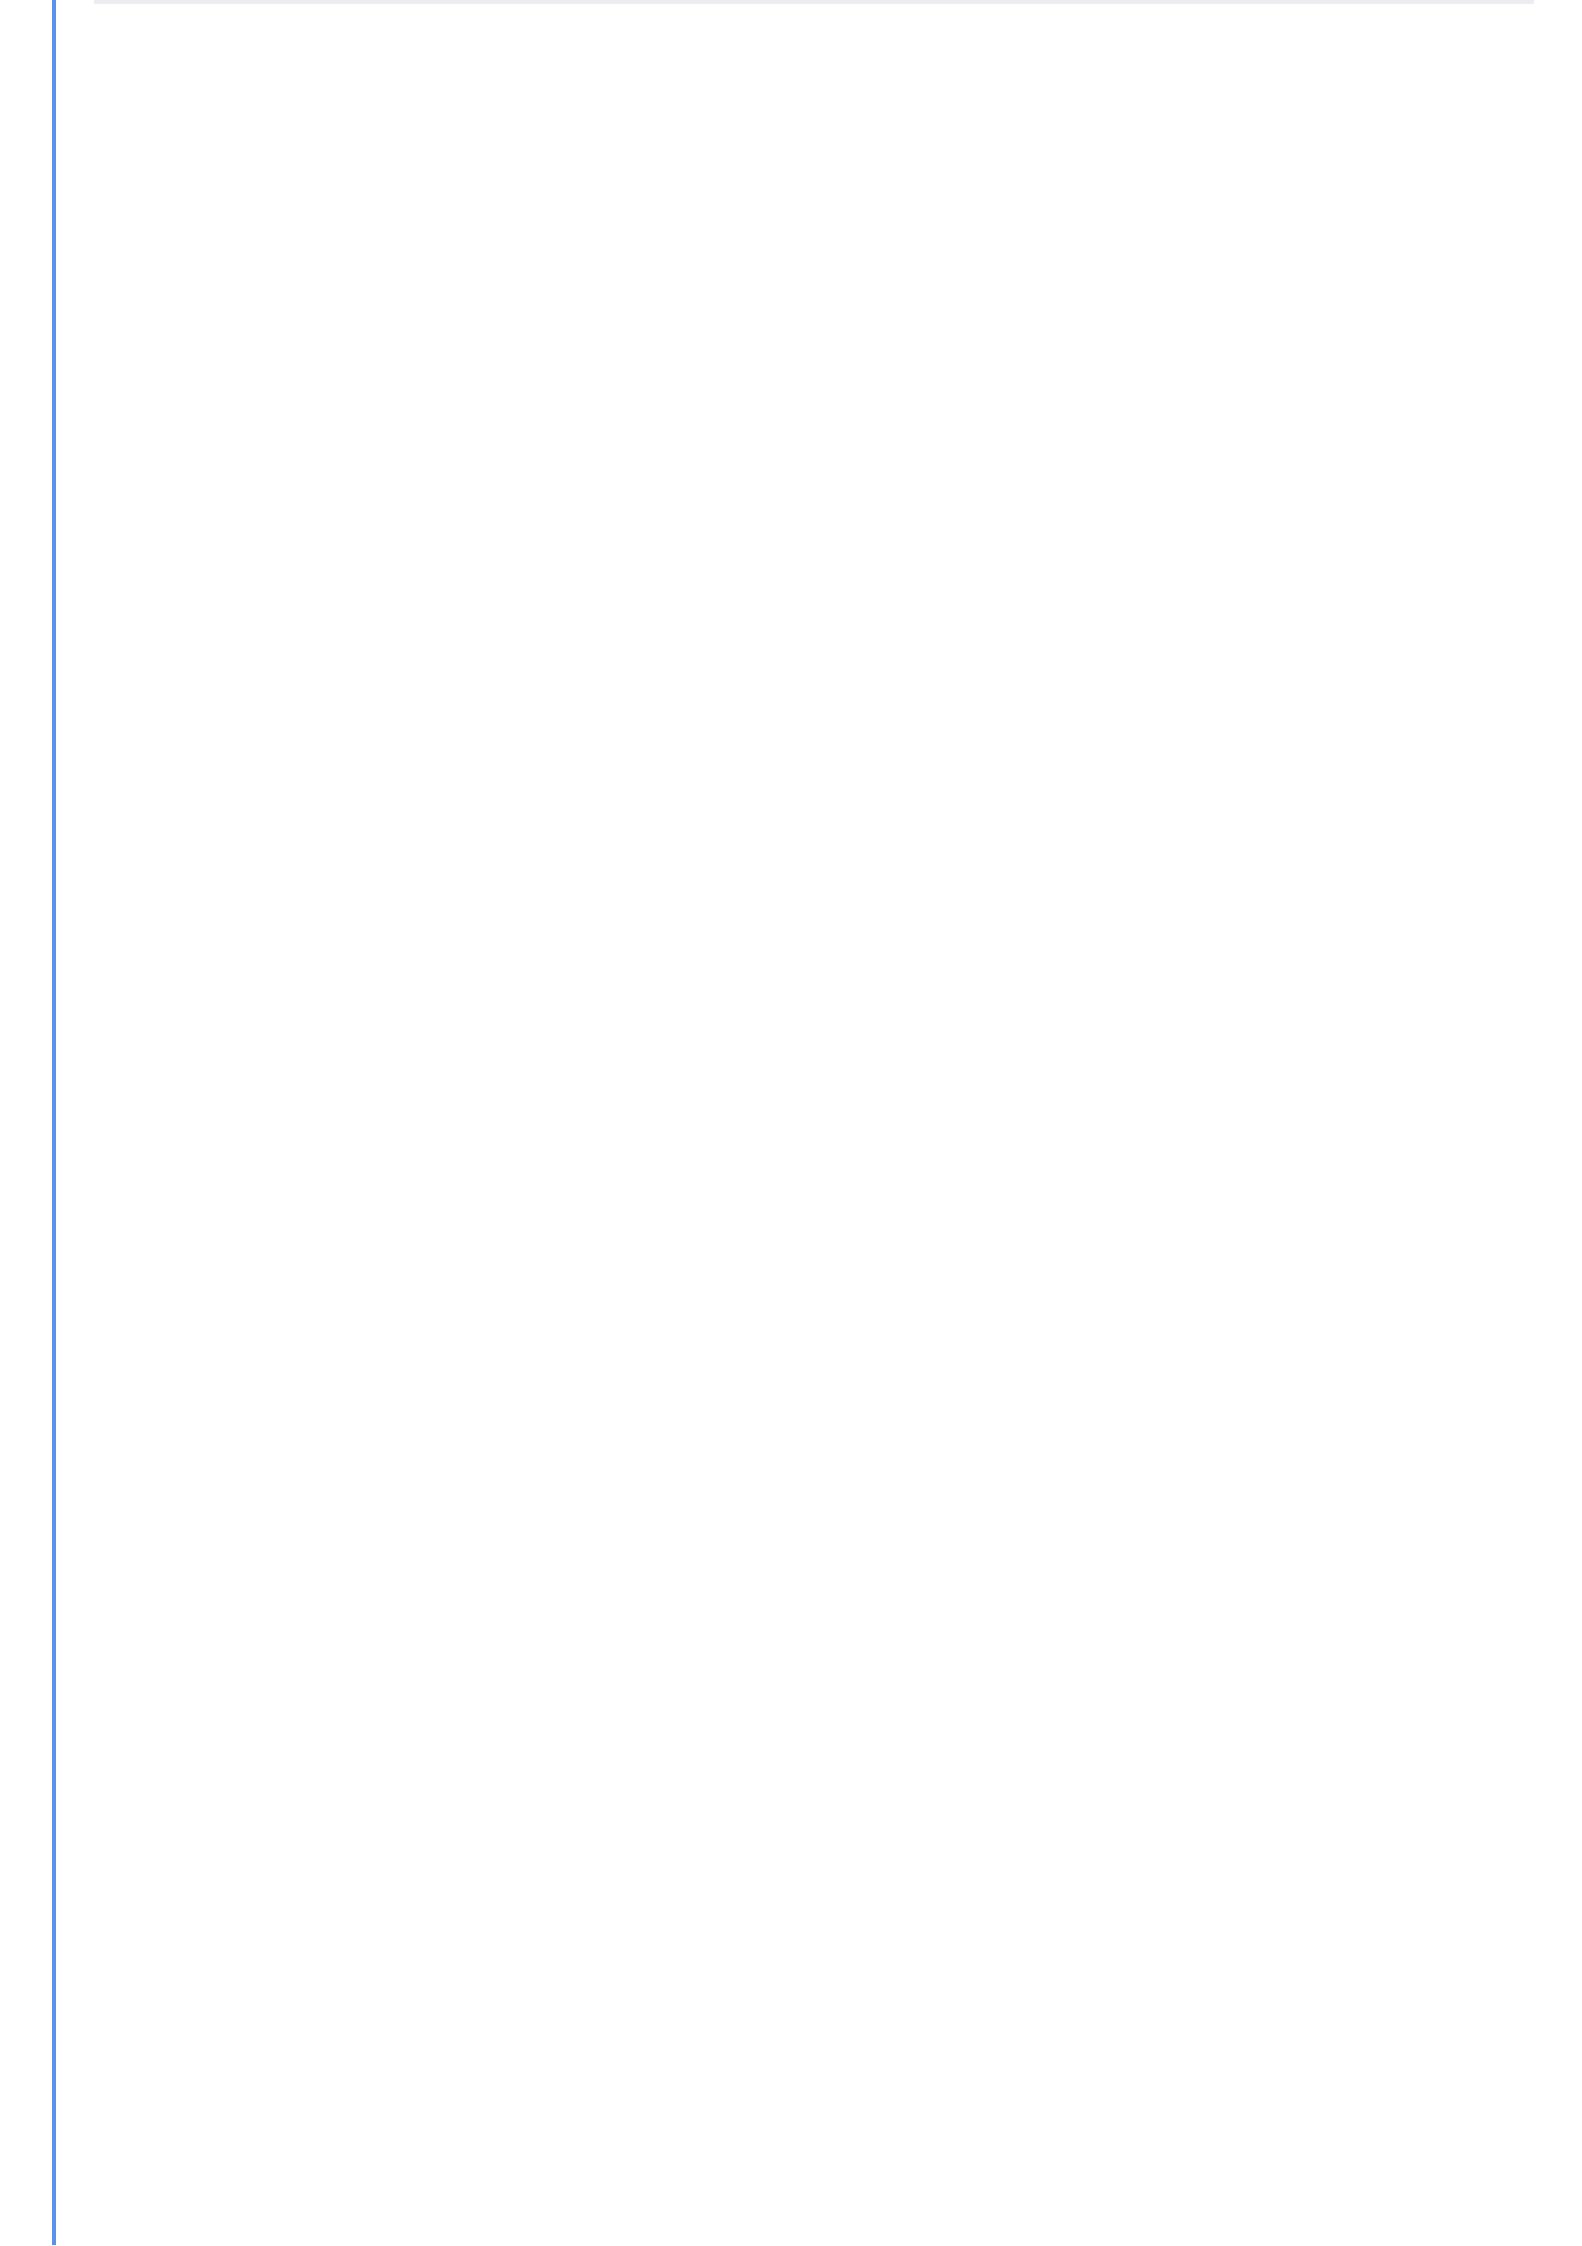
Digital Signatures

We use digital signatures as a way to show that a message came from a particular person (or holder of a key)

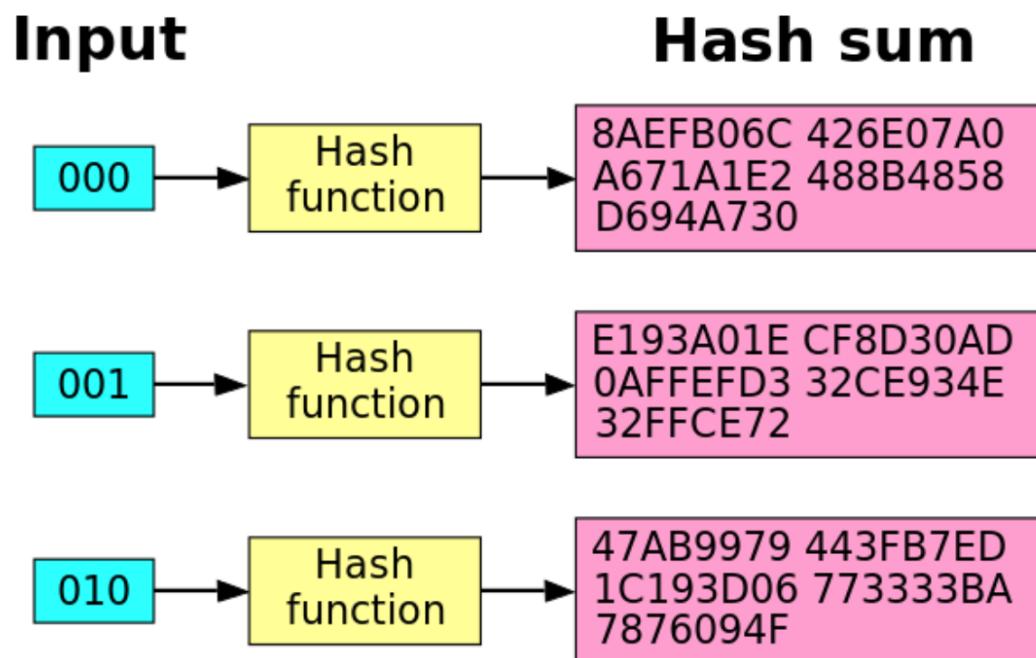
For example imagine Bob is signing a document and sending it to Alice.

A digital signature will have 4 properties

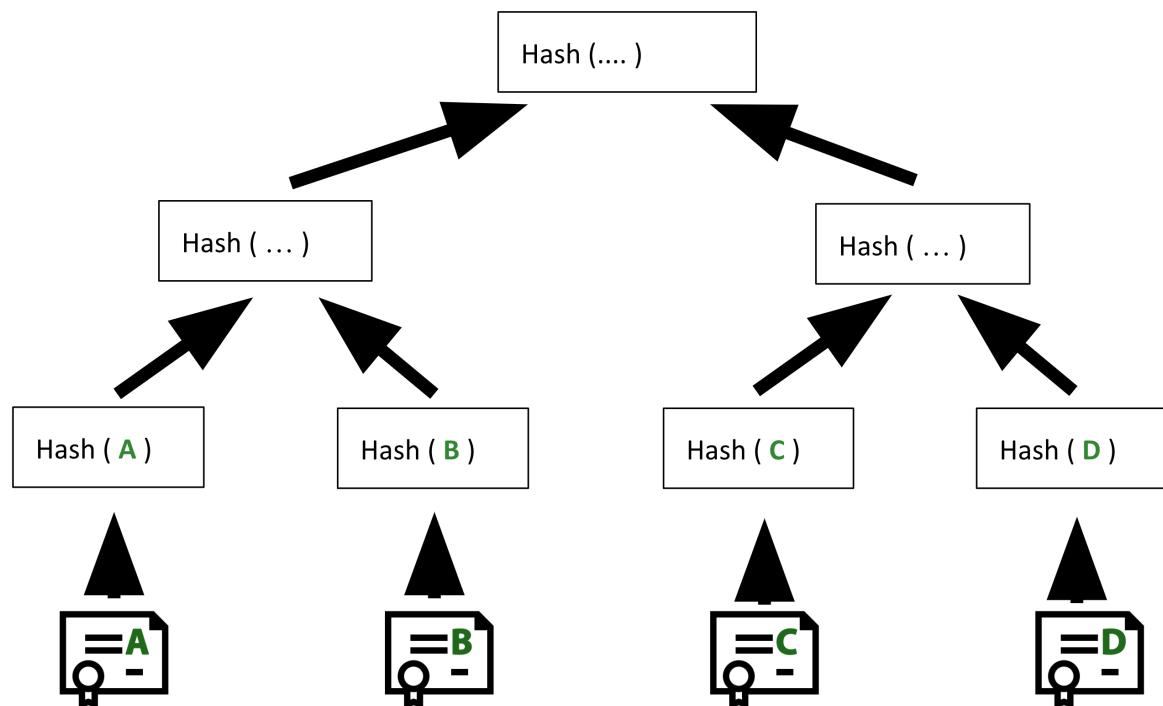
- Authenticity : when Alice verifies the message with Bob's public key, she knows that he signed the message.
- Unforgeable : only Bob knows his private key.
- Not reusable : the signature is tightly bound with the document, it cannot be transferred to any other document.
- Unalterable : If there is any alteration to the document, it will no longer be verifiable with Bob's public key.



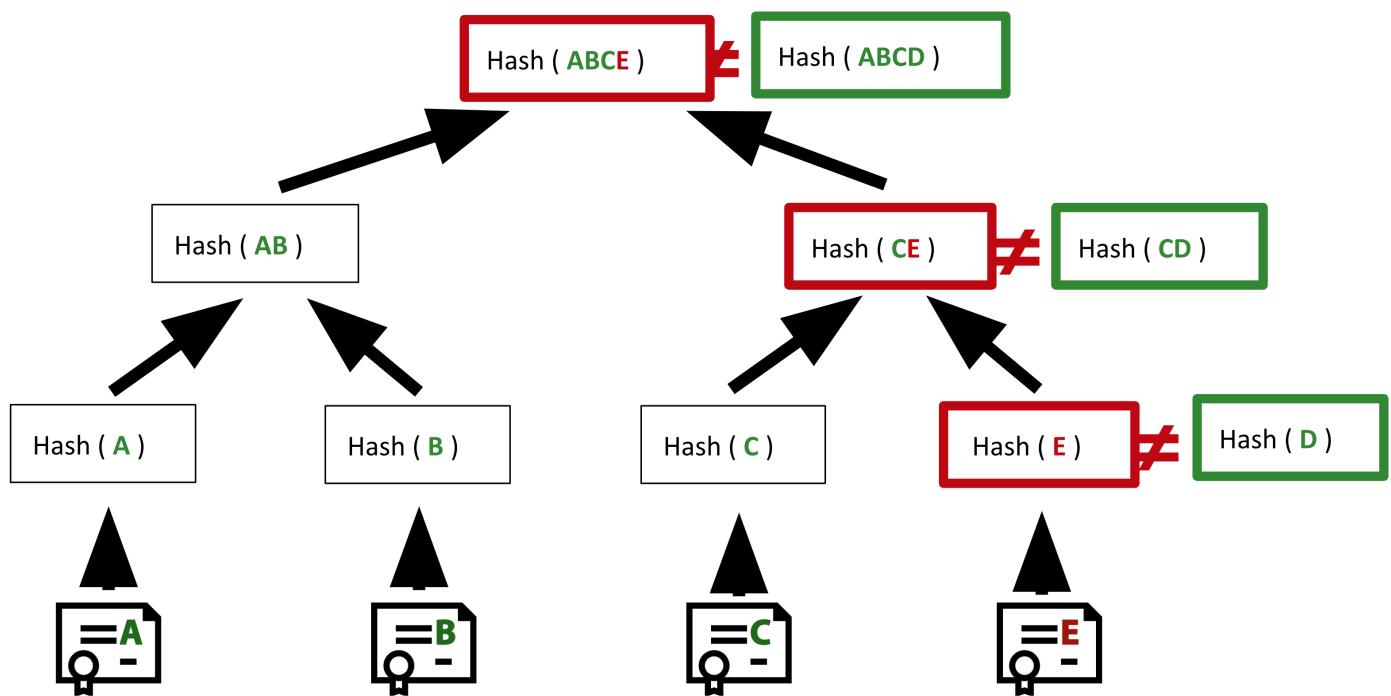
Cryptography - Hash Functions



Merkle Tree (the basic)



Merkle Tree (the basic)



That is the cryptographic background, how did people try to use this technology ?

The development of

- Electronic cash
 - Timestamping
 - P2P Systems
 - Consensus systems
-

1980s

David Chaum - Blind Signatures

David Chaum - DigiCash

1990s

Timestamping records

Adam Back - HashCash

Wei Dai - B-Money

2000s

Peer to peer networks

- Freenet / Gnutella / Bit Torrent

[Further Attempts at Electronic Cash](#)

"the one thing that's missing is a reliable e-cash, whereby on the internet you can transfer funds from A to B without A knowing B or B knowing A" - Milton Friedman 1999

1998 - b-money - Wei Dai (<http://www.weidai.com/bmoney.txt>)

1998 - Bit Gold - Nick Szabo (<https://nakamotoinstitute.org/bit-gold/>)

Bitcoin



Bitcoin QR Code



Satoshi Nakamoto is the name used by the presumed **pseudonymous** person or persons who developed **bitcoin**, authored the bitcoin **white paper**, and created and deployed bitcoin's original **reference implementation**



August 2008 - the domain name bitcoin.org registered

October 2008 - A Peer-to-Peer Electronic Cash System posted to a cryptography mailing list

January 2009 - Software implementation released as open source

2010, the first known commercial transaction using bitcoin occurred when programmer Laszlo Hanyecz bought two Papa John's pizzas for 10,000 BTC

[General Blockchain events since 2009](#)

2014 - Ethereum created

2017 - ICO Boom / Alternatives to Ethereum

2018 - Crypto winter

2020 - DeFi summer

2021 - Rise of NFTs / Gaming

2022 - Ethereum Merge

2022 - Another crypto winter

Layer1Theory

A simplistic / traditional view of blockchains

Typical Blockchain components

Gossip network



Shared public ledger



**

Cryptography



**

These components give the blockchain

- Transparency and verifiable state based on consensus
- Resilience

- Censorship resistance
 - Tamper proof interactions
-

Blockchain components (Bitcoin / Ethereum) in more detail

- A peer-to-peer (P2P) network connecting participants and propagating transactions and blocks of verified transactions, based on a standardised "gossip" protocol
- Messages, in the form of transactions, representing state transitions
- A set of consensus rules, governing what constitutes a transaction and what makes for a valid state transition
- A state machine that processes transactions according to the consensus rules
- A chain of cryptographically secured blocks that acts as a journal of all the verified and accepted state transitions
- A consensus algorithm that decentralizes control over the blockchain, by forcing participants to cooperate in the enforcement of the consensus rules
- A game-theoretically sound incentivization scheme (e.g., proof-of-work costs plus block rewards) to economically secure the state machine in an open environment
- One or more open source software implementations of the above ("clients")

[State Machine](#)

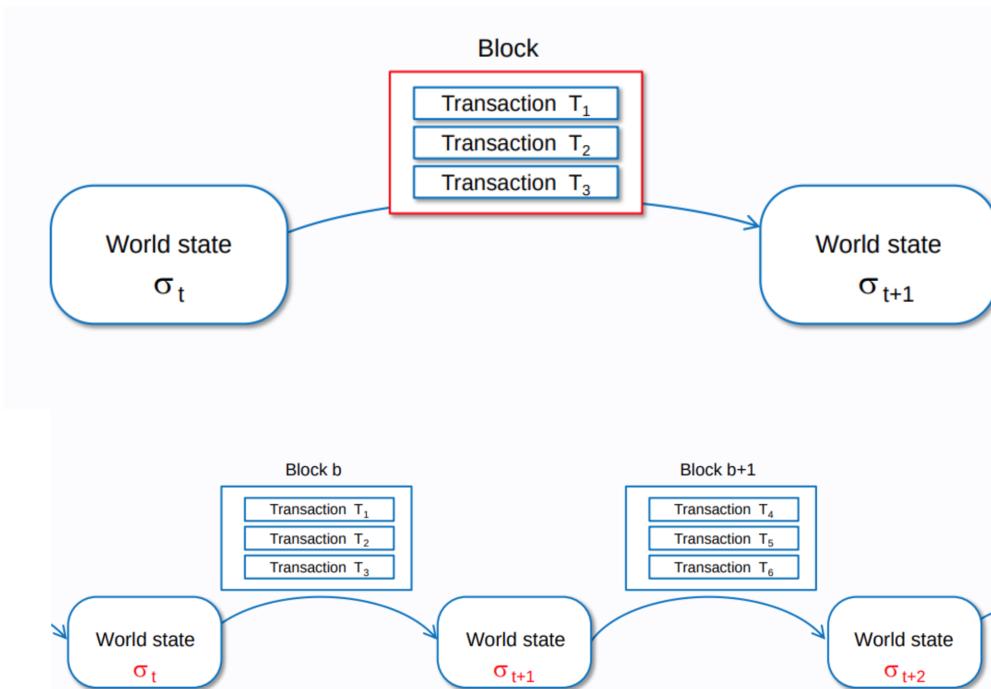
A useful way to think of the interaction between blocks, transactions and blockchain state is to model the blockchain as a state machine.

The state transition is caused by the execution of a set of transactions gathered in a block.

The use of a block is helpful to delineate the state transition, but is it not universally used.

Aptos / Sui and to some extend Solana focus more on the transactions.

Blockchain as a state machine



From : https://takenobu-hs.github.io/downloads/ethereum_evm_illustrated.pdf

Traditional Blockchain architecture

Network

From : [Bitcoin book](#)

Bitcoin is structured as a peer-to-peer network architecture on top of the internet. The term peer-to-peer, or P2P, means that the computers that participate in the network are peers to each other, that they are all equal, that there are no "special" nodes, and that all nodes share the burden of providing network services.

The network nodes interconnect in a mesh network with a "flat" topology. There is no server, no centralised service, and no hierarchy within the network.

Blockchain nodes

Nodes will typically

- Accept and transmit transactions (if valid)
- Keep a mempool of pending transactions
- Provide network discovery and routing functions
- Drop connections to misbehaving nodes

- Accept blocks and update their local ledger

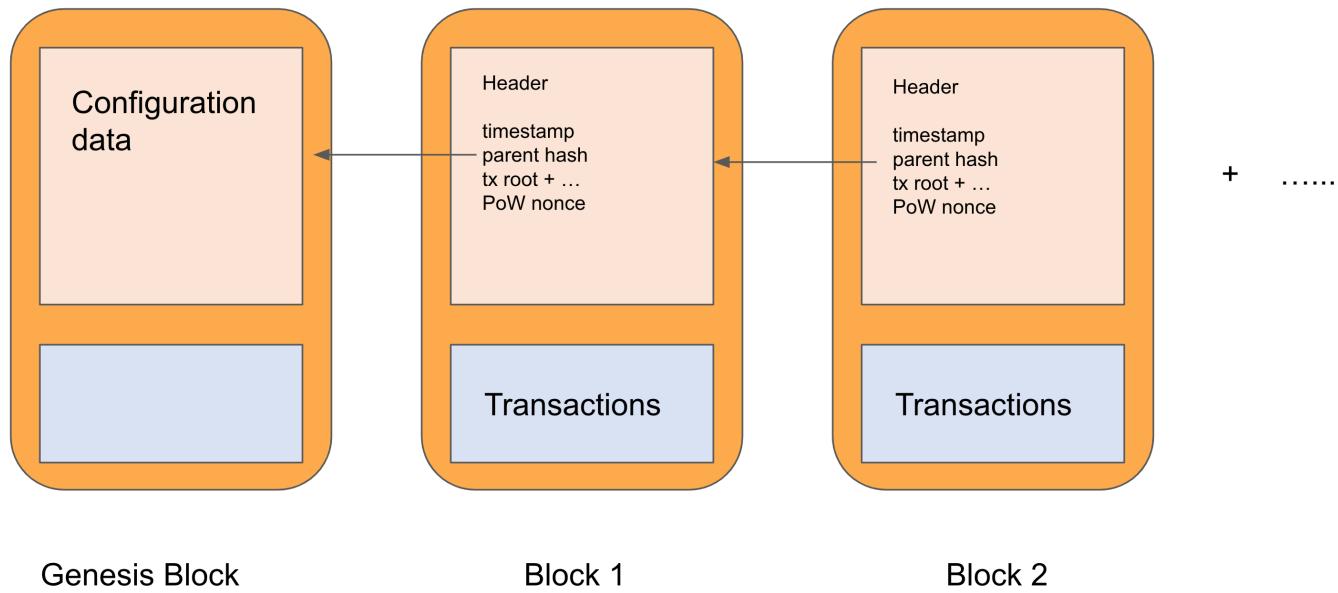
The connections are not based on geographical proximity but proximity in a hash table

Node discovery is via

- DNS seed nodes
- A locally stored list

[Block data structures](#)

Blockchain Data structure



Bitcoin Genesis Block

Raw Hex Version

00000000	01 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
00000010	00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
00000020	00 00 00 00 3B A3 ED FD 7A 7B 12 B2 7A C7 2C 3E;Íýz{.^zç,>
00000030	67 76 8F 61 7F C8 1B C3 88 8A 51 32 3A 9F B8 AA	gv.a.È.Ã^ŠQ2:Ý,^
00000040	4B 1E 5E 4A 29 AB 5F 49 FF FF 00 1D 1D AC 2B 7C	K.^J)«_Iýý...¬+
00000050	01 01 00 00 00 01 00 00 00 00 00 00 00 00 00 00
00000060	00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
00000070	00 00 00 00 00 00 FF FF FF FF 4D 04 FF FF 00 1DÝýýýM.Ýý..
00000080	01 04 45 54 68 65 20 54 69 6D 65 73 20 30 33 2F	..EThe Times 03/
00000090	4A 61 6E 2F 32 30 30 39 20 43 68 61 6E 63 65 6C	Jan/2009 Chancel
000000A0	6C 6F 72 20 6F 6E 20 62 72 69 6E 6B 20 6F 66 20	lor on brink of
000000B0	73 65 63 6F 6E 64 20 62 61 69 6C 6F 75 74 20 66	second bailout f
000000C0	6F 72 20 62 61 6E 6B 73 FF FF FF FF 01 00 F2 05	or banksÝýýý..ð.
000000D0	2A 01 00 00 00 43 41 04 67 8A FD B0 FE 55 48 27	*....CA.gŠý°þUH'
000000E0	19 67 F1 A6 71 30 B7 10 5C D6 A8 28 E0 39 09 A6	.gñ;q0..\\Ö"(à9.
000000F0	79 62 E0 EA 1F 61 DE B6 49 F6 BC 3F 4C EF 38 C4	ybaê.ap‰Iöê?Li8À
00000100	F3 55 04 E5 1E C1 12 DE 5C 38 4D F7 BA 0B 8D 57	óU.å.Á.þ\8M+ø..W
00000110	8A 4C 70 2B 6B F1 1D 5F AC 00 00 00 00	ŠLp+kñ._¬....

[By MikeG001 - Own work, CC BY-SA 4.0,](#)

Storage data structures

Merkle trees or variants such as Merkle Patricia trees are widely used to organise storage, though this can be an abstraction, the physical storage of items, including blocks and blockchain state is up to the design of the node software.

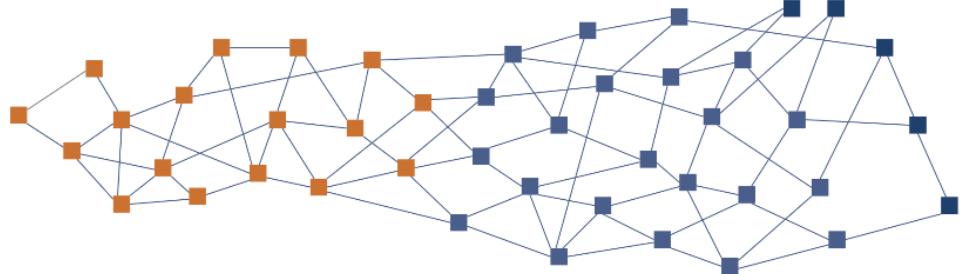
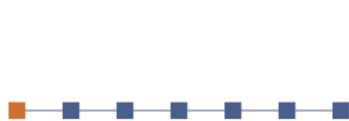
In Ethereum roadmap Merkle trees are planned to be replaced by Verkle trees for better performance.

Alternative data structures

Some chains such as Sui and IOTA use a Directed Acyclic Graph (DAG) rather than a linear list of blocks.

DAGs are a structure where there are no loops, and the items are linked with directed links.

Traditional blockchains can be seen as a specific case of a DAG in which each block only has one parent, DAGs in general allow multiple parents.



Modular Blockchains

Execution



Fuel



Optimism



polygon Hermez



Scroll



STARKNET



zkSync



ARBITRUM



Aztec

Settlement/Consensus



ethereum



sui



APTES



CELESTIA



COSMOS



Evmos

Data Availability



zkPorter



CELESTIA



ethereum



STARKEx

Eigenlayer/Datalayer



If we follow the principle of separation of concerns, we can use a combination of blockchains to provide the functionality of a L1 and increase scalability.

For further details see Volt [article](#)

The Blockspace Market

See [article](#)

Demand Side

The demand side is formed from the collection of transactions supplied by users (via DApps)

Supply Side

The supply side is provided by the validators proposing and voting on blocks, adding to the Ethereum blockchain.

The 2 sides are mediated by congestion and fees.

The mechanics are complex and give rise to features such as

- MEV
- ASIC development

We can see from this, that the functionality we need to provide falls into 3 categories

- Execution
- Settlement / Consensus
- Data availability

Execution

This is where smart contracts or programs implement the business logic.

We can design the environment to optimise the execution, which means we can move away from the EVM if desired.

Some execution layers such as Starknet use a completely different execution environment, and smart contract language to Ethereum L1.

Settlement Consensus

In this layer we come to agreement about the state transitions, using consensus mechanisms for security and decentralisation.

The execution layer will send updates to the settlement layer, which once accepted are seen as securing the updates to the execution layer.

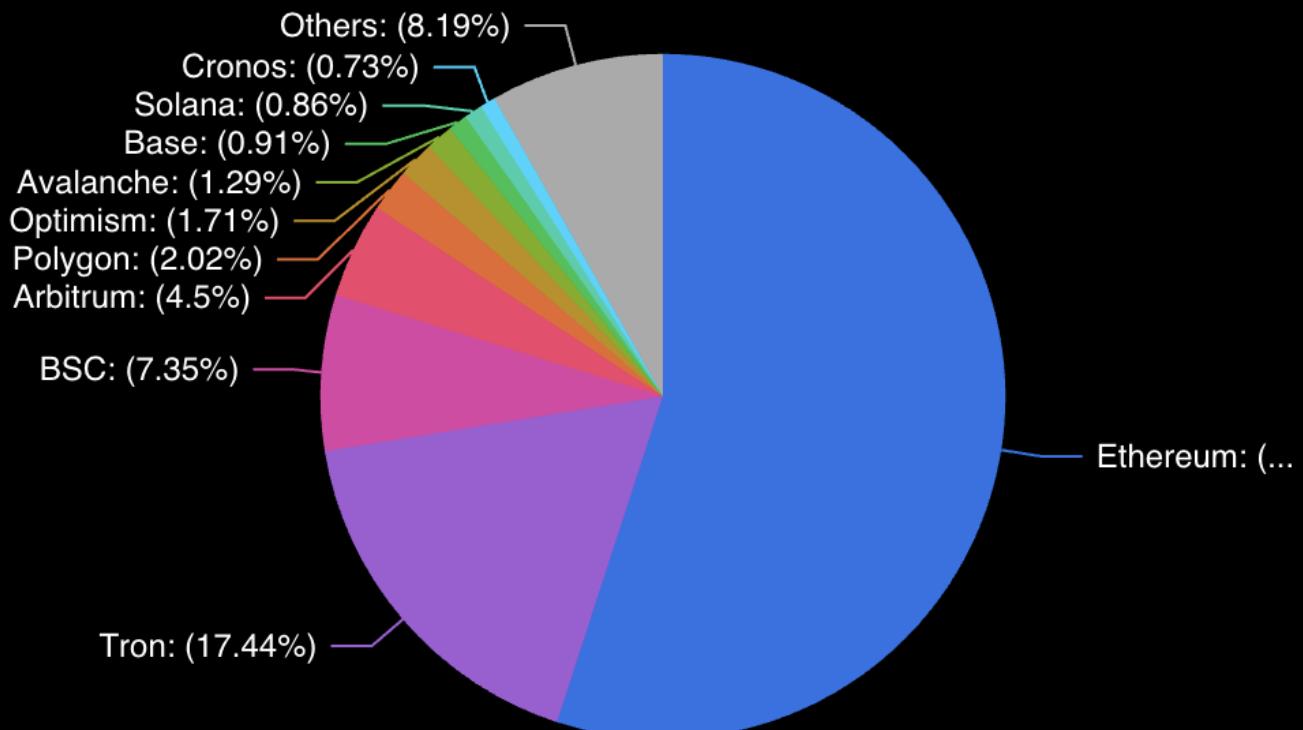
Data Availability

Data availability is the ability to verify that the information for a proposed block was actually published on the blockchain network. A block is "unavailable" if the producer publishes the block itself but withholds data used to create the block.

Chains overview

From Defi Llama

Total Value Locked All Chains



Consensus on Ethereum

From Ethereum developer [documentation](#)

"Now technically, proof-of-work and proof-of-stake are not consensus protocols by themselves, but they are often referred to as such for simplicity. They are actually Sybil resistance mechanisms and block author selectors; they are a way to decide who is the author of the latest block. It is this Sybil resistance mechanism combined with a chain selection rule that makes up a true consensus mechanism."

There are 2 parts to block addition :

- block producer selection
- block acceptance

From yellow paper

" Since the system is decentralised and all parties have an opportunity to create a new block on some older pre-existing block, the resultant structure is necessarily a tree of blocks. In order to form a consensus as to which path, from root (the genesis block) to leaf (the block containing the most recent transactions) through this tree structure, known as the blockchain, there must be an agreed-upon scheme. If there is ever a disagreement between nodes as to which root-to-leaf path down the block tree is the 'best' blockchain, then a fork occurs."

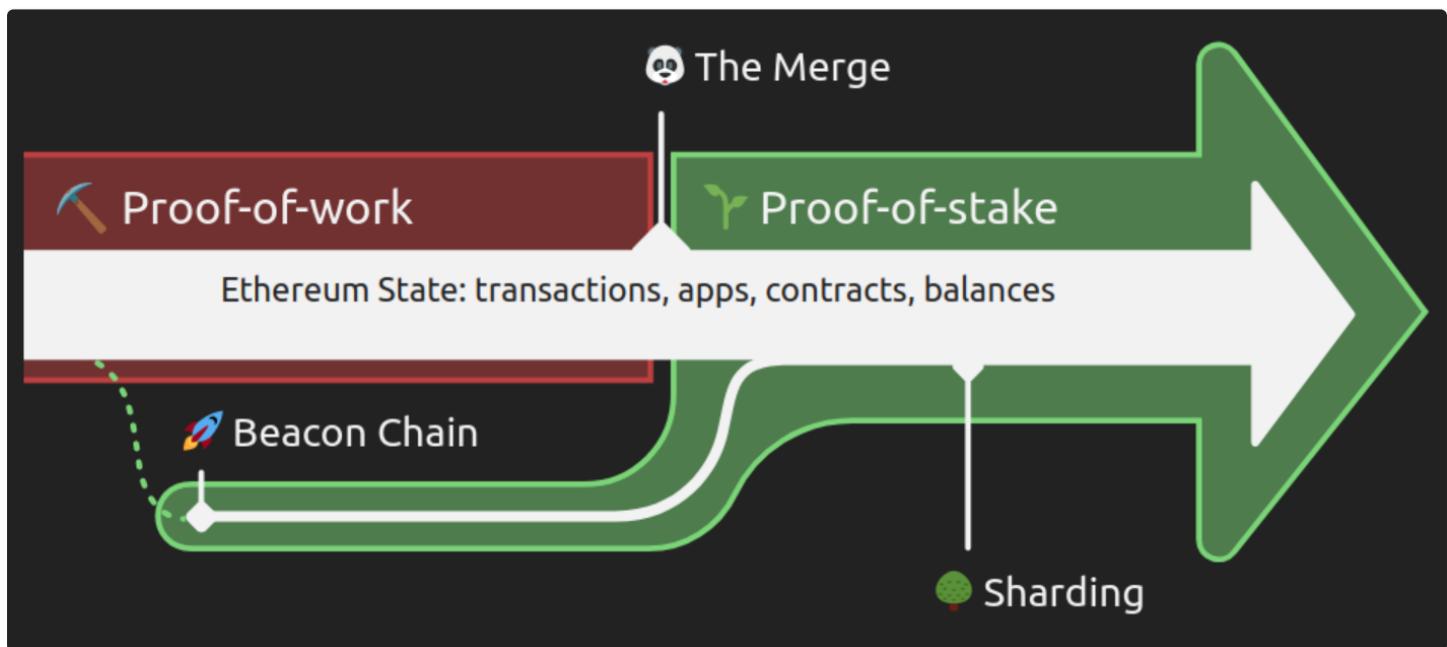
[Consensus Mechanism Families](#)

- Nakamoto style - Bitcoin, Ethereum pre merge
 - BFT style - Ethereum post merge
 - Avalanche Snow
-

The Merge - Proof of stake update

- Replaced proof of work with the proof of stake beacon chain.
i.e. merging existing beacon chain into ethereum.
- Pre merge the Beacon Chain had not been processing Mainnet transactions. Instead, it has been reaching consensus on its own state by agreeing on active validators and their account balances.
- POS specs: <https://github.com/ethereum/consensus-specs#phase-0>

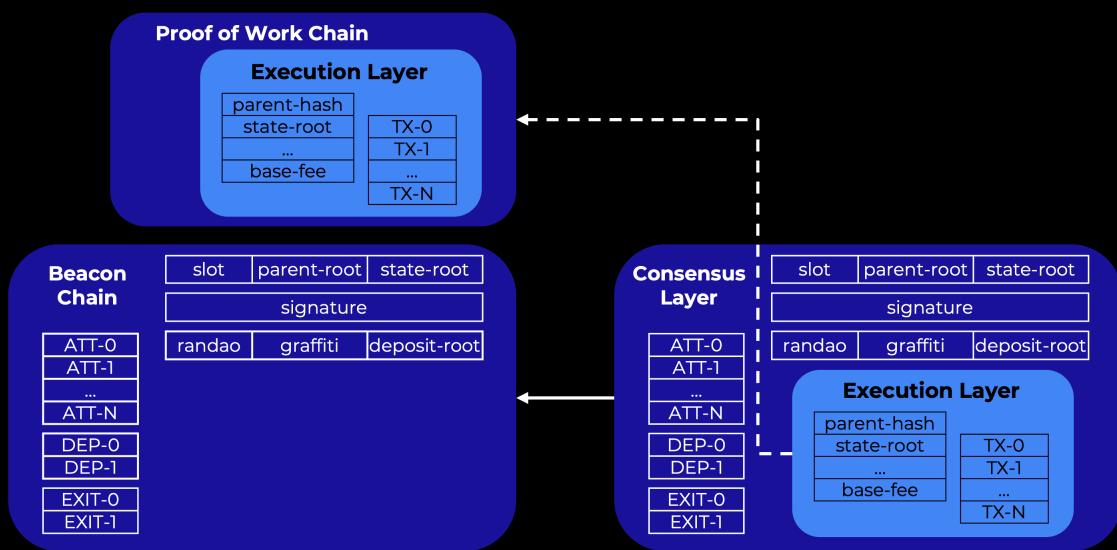
"The number 587500000000000000000000 joins the list of the most important integers of our time. The moment Ethereum's proof of work chain accumulates that much difficulty (read, hashes done), the entire network will switch over to proof of stake"





The Merge

Plugging the Execution Layer into the Beacon Chain Consensus Layer

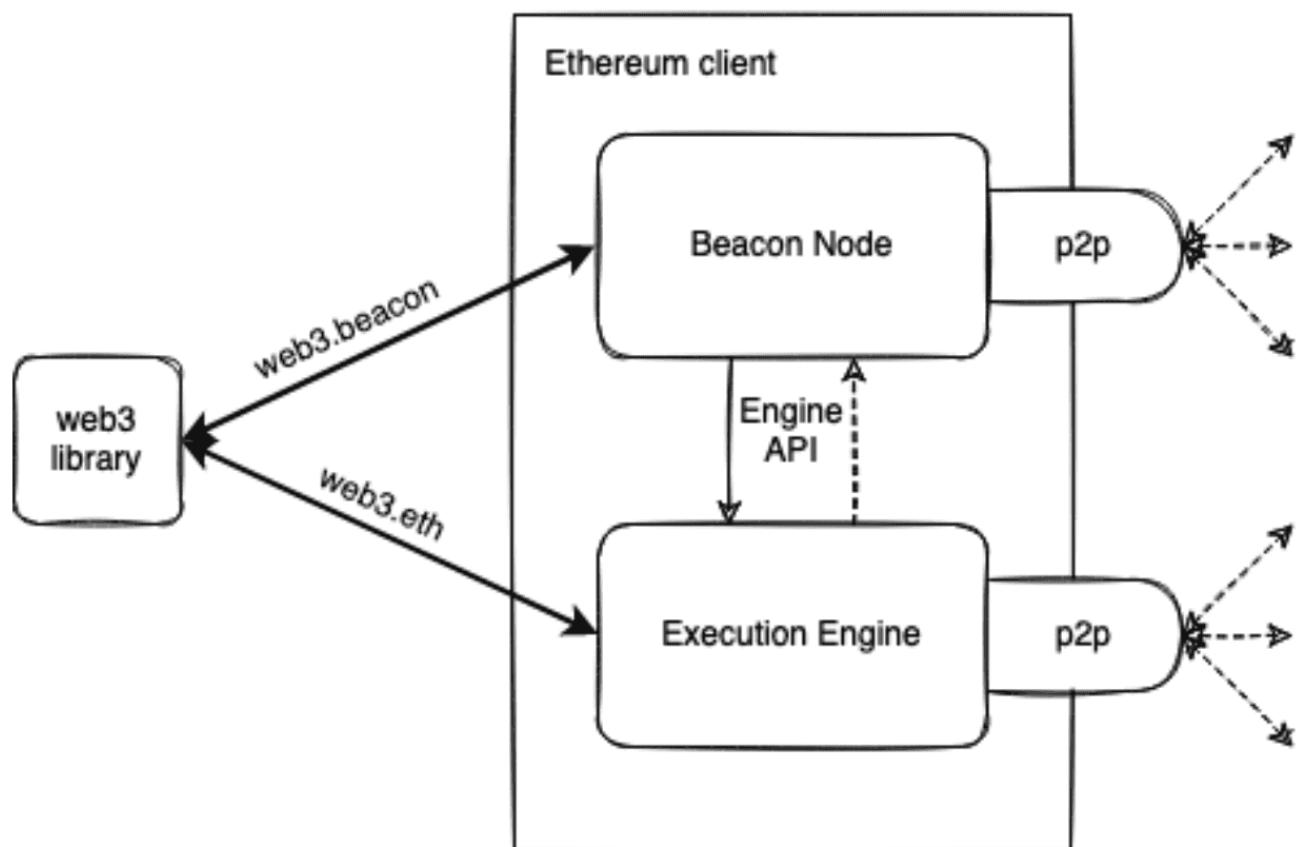


Source: Tim Beiko, Ethereum Foundation



Ethereum clients after the merge

- Eth 1.0 clients continue to handle execution. They process blocks, maintain mempools, and manage and sync state. The PoW stuff has been taken out.
- Consensus client – Current Beacon Chain clients continue to handle PoS consensus. They track the chain's head, gossip and attest to blocks, and receive validator rewards.



Execution Clients

See [Docs](#)

Client	Language	Operating Systems	Networks
Geth ↗	Go	Linux, Windows, macOS	Mainnet, Sepolia, Goerli
Nethermind ↗	C#, .NET	Linux, Windows, macOS	Mainnet, Sepolia, Goerli, and more
Besu ↗	Java	Linux, Windows, macOS	Mainnet, Sepolia, Goerli, and more
Erigon ↗	Go	Linux, Windows, macOS	Mainnet, Sepolia, Goerli, and more
Reth ↗	Rust	Linux, Windows, macOS	Mainnet, Sepolia, Goerli, and more

Consensus Clients

See [Docs](#)

Client	Language	Operating systems	Networks
Lighthouse ↗	Rust	Linux, Windows, macOS	Beacon Chain, Goerli, Pyrmont, Sepolia, Ropsten, and more
Lodestar ↗	TypeScript	Linux, Windows, macOS	Beacon Chain, Goerli, Sepolia, Ropsten, and more
Nimbus ↗	Nim	Linux, Windows, macOS	Beacon Chain, Goerli, Sepolia, Ropsten, and more
Prysm ↗	Go	Linux, Windows, macOS	Beacon Chain, Gnosis, Goerli, Pyrmont, Sepolia, Ropsten, and more
Teku ↗	Java	Linux, Windows, macOS	Beacon Chain, Gnosis, Goerli, Sepolia, Ropsten, and more

A number of different synchronisation modes are available see [docs](#)

[Client diversity](#) in the network is very much encouraged

Consensus after the Merge

Ethereum has moved to [Gasper](#) (Casper FFG + LMD GHOST (Latest Message Driven Greediest Heaviest Observed SubTree))

Consensus relies on both LMD-GHOST – which adds new blocks and decides what the head of the chain is – and Casper FFG which makes the final decision on which blocks *are* and *are not* a part of the chain. GHOST's favourable liveness properties allow new blocks to quickly and efficiently be added to the chain, while FFG follows behind to provide safety by finalising epochs.

The two protocols are merged by running GHOST from the last finalised block as decided upon by FFG.

By construction, the last finalised block is always a part of the chain which means GHOST doesn't need to consider earlier blocks.

Safety favouring protocols such as Tendermint can halt, if they don't get enough votes.

Liveness favouring protocols such as Nakamoto continue to add blocks, but they may not come to finality.

Ethereum will achieve finality by checkpointing

[Epochs and checkpointing](#)

Epochs of about 6 mins have 32 slots with all validators attesting to one slot (~12K attestations per block)

The fork-choice rule LMD GHOST then determines the current head of chain based on these attestations.

Finality is achieved when sufficient votes are reached, generally after 2 epochs.

Validator Selection and consensus in more detail

You can find stats about the validators [here](#)

⌚ Most recent epochs					🎲 Most recent blocks			
Epoch	Time	Final	Eligible (ETH)	Voted	Epoch	Slot	Block	Status
178,858	3 mins ago	No	Calculating...	Calculating...	178,858	5,723,471	16,554,754	Proposed
178,857	9 mins ago	No	16,379,848	15,807,513 (96.51%)	178,858	5,723,470	16,554,753	Proposed
178,856	16 mins ago	Yes	16,379,624	16,316,025 (99.61%)	178,858	5,723,469	16,554,752	Proposed
178,855	22 mins ago	Yes	16,379,400	16,304,281 (99.54%)	178,858	5,723,468	16,554,751	Proposed
178,854	28 mins ago	Yes	16,379,400	16,323,577 (99.66%)	178,858	5,723,467	16,554,750	Proposed
178,853	35 mins ago	Yes	16,379,400	15,816,985 (96.57%)	178,858	5,723,466	16,554,749	Proposed
178,852	41 mins ago	Yes	16,379,400	16,312,281 (99.59%)	178,858	5,723,465	16,554,748	Proposed
178,851	48 mins ago	Yes	16,379,400	16,321,497 (99.65%)	178,858	5,723,464	16,554,747	Proposed
178,850	54 mins ago	Yes	16,379,400	16,280,569 (99.40%)	178,858	5,723,463	16,554,746	Proposed
178,849	1 hr ago	Yes	16,379,400	16,323,225 (99.66%)	178,858	5,723,462	16,554,745	Proposed

A slot occurs every 12s and one validator is chosen to submit a block within that slot.

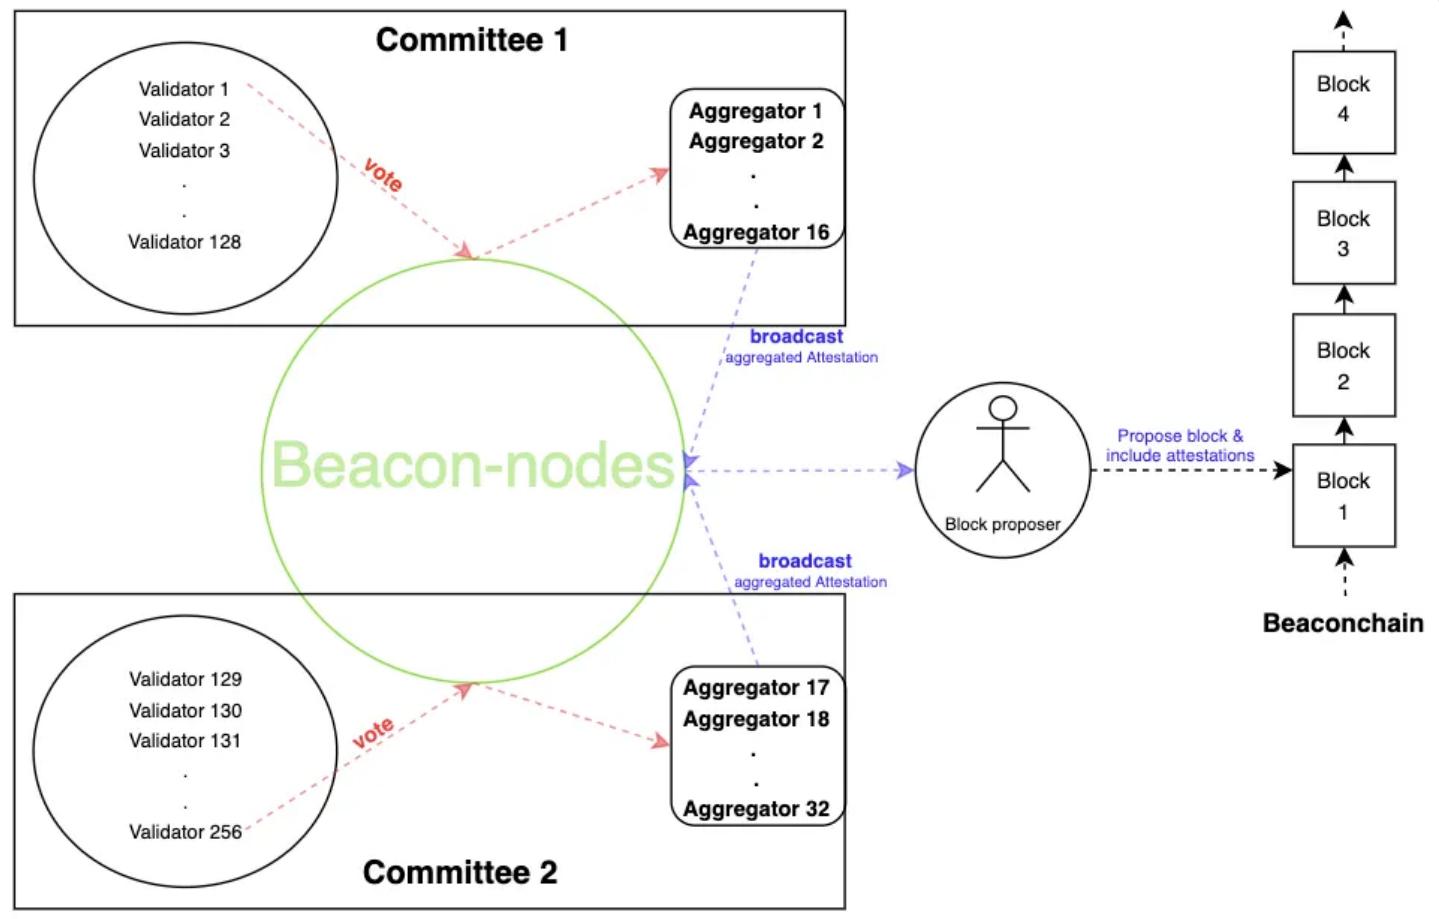
If the validator fails to do so, the slot is let empty.

The first block within an epoch (~6.4 mins) is a checkpoint block.

Coming to consensus about the block

If a validator is not chosen to produce a block, it will instead vote on what it regards as the current head of the chain and the checkpoint block.

Within an epoch a validator will only vote once.



Validators are grouped into committees at random, their votes are aggregated and published in the block header.

[This](#) article gives an in depth view of validator rewards

Fork choice rule

When faced with a potential fork, we choose the fork that has the most votes, but when counting the votes, we only include the last one from any validator.

[Voting rules](#)

The validator must vote for the chain they consider to be correct
The validator cannot vote for 2 blocks in any one time slot.

Finalisation

Validators vote on a pair of checkpoint blocks, to indicate that they are valid.

Once a checkpoint block gets sufficient votes, it is regarded as 'justified'
Once its child checkpoint block becomes justified, then the parent is regarded as final.